

Scanned Page Sample

Paper documents have carried information for centuries, but the last three decades have pushed most of that record into digital form. Converting a scanned page or a legacy file into structured text is a small task in isolation, yet it sits underneath search engines, archives, and every pipeline that needs to reason about what a document actually says rather than just where it is stored.

A text extraction tool is judged by more than whether it runs without crashing. It has to preserve reading order across columns, recognize headings as headings rather than as ordinary paragraphs, and produce output that a downstream system can actually use. Many tools clear the first bar and fail the second, which is why word count alone is a weak signal of quality.